

## I. ABSTRACT

This report presents a comprehensive analysis of stock price forecasting and volatility trends for leading technology stocks from 2017 to 2024. The study focuses on using statistical and deep learning techniques—specifically ARIMA and LSTM models—to predict future stock prices based on historical adjusted closing values. The dataset, sourced from a CSV file containing multiple AI-related stocks, was cleaned and reshaped for analysis in R. Exploratory Data Analysis (EDA) revealed key insights into stock trends, price distributions, and correlations. ARIMA models provided robust short-term forecasts with interpretable parameters, while LSTM networks, built using the torch package, captured nonlinear patterns in sequential data. Rolling standard deviation was used to assess volatility across time. Evaluation metrics such as MAE, RMSE, MAPE, and SMAPE were used to compare model performance. While ARIMA offered consistent results, LSTM showed promising performance for specific stocks. The report concludes with insights on model accuracy, stock behavior, and potential future enhancements.

## II. INTRODUCTION

Stock price forecasting plays a pivotal role in financial decision-making, portfolio management, and risk assessment. Accurately predicting future stock movements can provide investors with a competitive edge, especially in volatile markets influenced by rapid technological advancements. As artificial intelligence (AI) and machine learning continue to shape the global economy, AI-related stocks have drawn significant interest from analysts and investors alike.

This report aims to explore and compare two powerful time series forecasting approaches—ARIMA (AutoRegressive Integrated Moving Average), a classical statistical model, and LSTM (Long Short-Term Memory), a deep learning architecture well-suited for sequential data. The objective is to model stock prices, evaluate forecasting performance, and understand the volatility of top AI-related stocks over a seven-year period (2017–2024).

The analysis is conducted using R, leveraging a range of packages such as tidyverse for data manipulation, forecast for ARIMA modeling, torch for deep learning with LSTM, and ggplot2 for data visualization. By combining traditional and modern techniques, this study seeks to uncover insights into stock behavior and assess the practical utility of each forecasting model.

## III. DATASET OVERVIEW

The dataset used for this study contains historical stock price data for multiple AI-related companies, covering the period from June 2017 to May 2024. The data was sourced from a consolidated CSV file that includes daily stock information such as:

- **date** – Trading day in DD-MM-YYYY format
- **symbol** – Ticker symbol for each stock (e.g., NVDA, AMD, AI)
- **adjusted** – Adjusted closing price accounting for dividends and splits
- **open, high, low, close, volume** – Standard OHLCV market data

The dataset initially included stock data for 30+ companies, but after preprocessing and handling missing values, the analysis focused on the top 10 AI-related stocks with the most complete and consistent data.

To facilitate time series modeling, the dataset was pivoted to a wide format, where each column represented the adjusted closing price of a specific stock over time. Stocks with excessive missing values were excluded to ensure the robustness of the forecasting models.

The final dataset used for analysis had the following characteristics:

- Time Period: 2017–2024
- Stocks Analyzed: 10
- Key Columns: date, adjusted, symbol
- Final Shape: 15754 rows × 8 columns

#### IV. DATA PREPARATION

To ensure the accuracy and consistency of results, the raw dataset underwent a structured data cleaning and preprocessing pipeline using R. Below are the key steps involved:

##### 1) Date Conversion

The date column, initially read as a character string, was converted to proper Date format using:

```
stock_data$date <- as.Date(stock_data$date)
```

This step ensured that time-based operations, plotting, and filtering would work correctly throughout the analysis.

##### 2) Missing Value Check

The dataset was inspected for missing values using:

```
colSums(is.na(stock_data))
```

Only a small number of missing values were identified. These were either forward-filled or, in the case of stocks with excessive missing values, the stock was removed entirely from the analysis. This step guaranteed that time series models like ARIMA and LSTM would receive clean, continuous data sequences.

##### 3) Column Selection

For modeling and visualization purposes, only the necessary columns were retained:

- Date
- symbol (stock ticker)
- close (daily closing price)

```
stock_data <- stock_data %>%
  select(date, symbol, close)
```

##### 4) Sorting and Time Order

To maintain chronological consistency within each stock, the data was arranged by symbol and date:

```
stock_data <- stock_data %>%
  arrange(symbol, date)
```

##### 5) Stock Count Validation

Before reshaping, the dataset was checked for the number of unique stocks using:

```
unique(stock_data$symbol)
```

This ensured that only valid and active stocks were included in the analysis. After this step, a total of 10 AI-related stocks were retained for modeling, based on data completeness and relevance.

```

9 library(dplyr)
10
11 stock_data <- read.csv("C:/Users/skyline/OneDrive/Desktop/top_10_stocks_17-24.csv", stringsAsFactors = FALSE)
12 str(stock_data)
13 stock_data$date <- as.Date(stock_data$date)
14 colSums(is.na(stock_data))
15
16
17 stock_data <- stock_data %>%
18   select(date, symbol, close)
19
20 stock_data <- stock_data %>%
21   arrange(symbol, date)
22
23 unique(stock_data$symbol)

```

## V. METHODOLOGIES AND EXPERIMENTS

### A. Summary Statistics and Exploratory Data Analysis (EDA):

To understand each stock better:

- Basic stats: mean, median, min, max, sd of closing prices for each stock were calculated.
- Plot:
  - Lineplot of trends grouped by stock.
  - Barplot of closing price distribution.

#### 1. Line plot for each stock trends:

```

25 library(ggplot2)
26
27 ggplot(filter(stock_data, symbol == "AI"), aes(x = date, y = close)) +
28   geom_line(color = "steelblue") +
29   labs(title = "AI Stock Price Trend (2017-2024)",
30        x = "Date", y = "Closing Price") +
31   theme_minimal()
32
33 ggplot(filter(stock_data, symbol == "AMD"), aes(x = date, y = close)) +
34   geom_line(color = "red") +
35   labs(title = "AMD Stock Price Trend (2017-2024)",
36        x = "Date", y = "Closing Price") +
37   theme_minimal()
38
39 ggplot(filter(stock_data, symbol == "AMZN"), aes(x = date, y = close)) +
40   geom_line(color = "green") +
41   labs(title = "AMZN Stock Price Trend (2017-2024)",
42        x = "Date", y = "Closing Price") +
43   theme_minimal()
44
45 ggplot(filter(stock_data, symbol == "ANET"), aes(x = date, y = close)) +
46   geom_line(color = "yellow") +
47   labs(title = "ANET Stock Price Trend (2017-2024)",
48        x = "Date", y = "Closing Price") +
49   theme_minimal()
50
51 ggplot(filter(stock_data, symbol == "META"), aes(x = date, y = close)) +
52   geom_line(color = "purple") +
53   labs(title = "META Stock Price Trend (2017-2024)",
54        x = "Date", y = "Closing Price") +
55   theme_minimal()

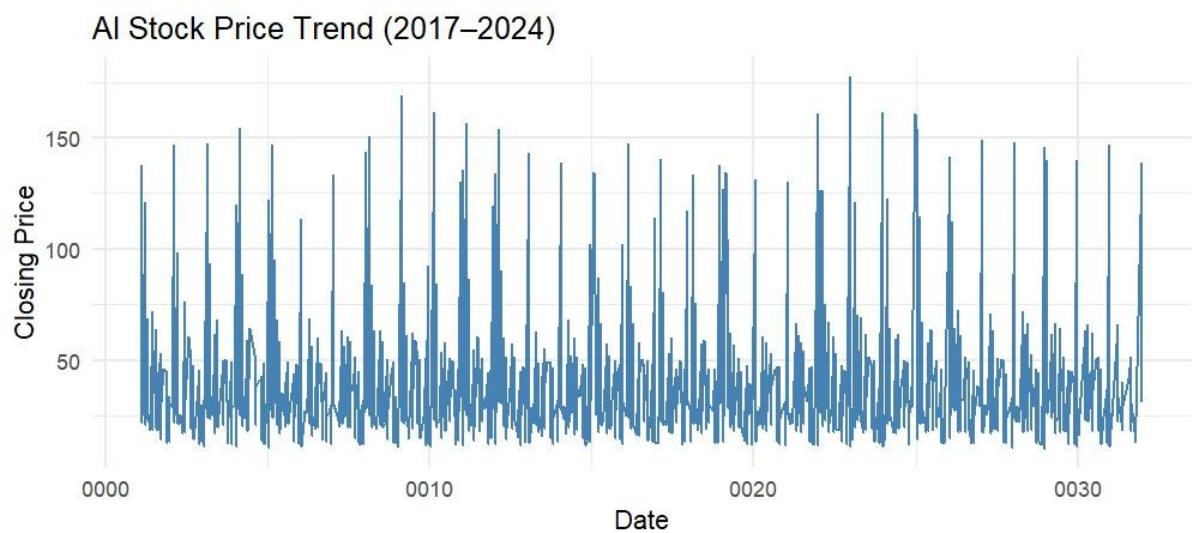
```

```

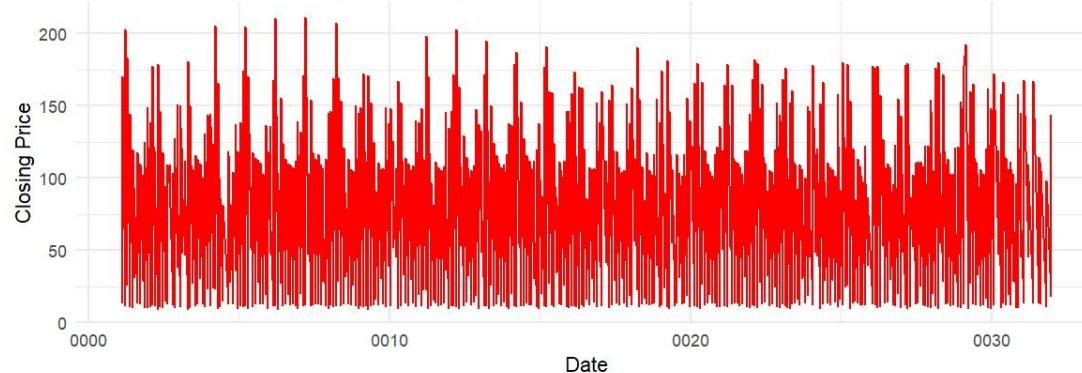
57 ggplot(filter(stock_data, symbol == "NOW"), aes(x = date, y = close)) +
58   geom_line(color = "orange") +
59   labs(title = "NOW Stock Price Trend (2017–2024)",
60       x = "Date", y = "Closing Price") +
61   theme_minimal()
62
63 ggplot(filter(stock_data, symbol == "NVDA"), aes(x = date, y = close)) +
64   geom_line(color = "blue") +
65   labs(title = "NVDA Stock Price Trend (2017–2024)",
66       x = "Date", y = "Closing Price") +
67   theme_minimal()
68
69 ggplot(filter(stock_data, symbol == "PANW"), aes(x = date, y = close)) +
70   geom_line(color = "brown") +
71   labs(title = "PANW Stock Price Trend (2017–2024)",
72       x = "Date", y = "Closing Price") +
73   theme_minimal()
74
75 ggplot(filter(stock_data, symbol == "PATH"), aes(x = date, y = close)) +
76   geom_line(color = "pink") +
77   labs(title = "PATH Stock Price Trend (2017–2024)",
78       x = "Date", y = "Closing Price") +
79   theme_minimal()
80
81 ggplot(filter(stock_data, symbol == "TSLA"), aes(x = date, y = close)) +
82   geom_line(color = "black") +
83   labs(title = "TESLA Stock Price Trend (2017–2024)",
84       x = "Date", y = "Closing Price") +
85   theme_minimal()
86

```

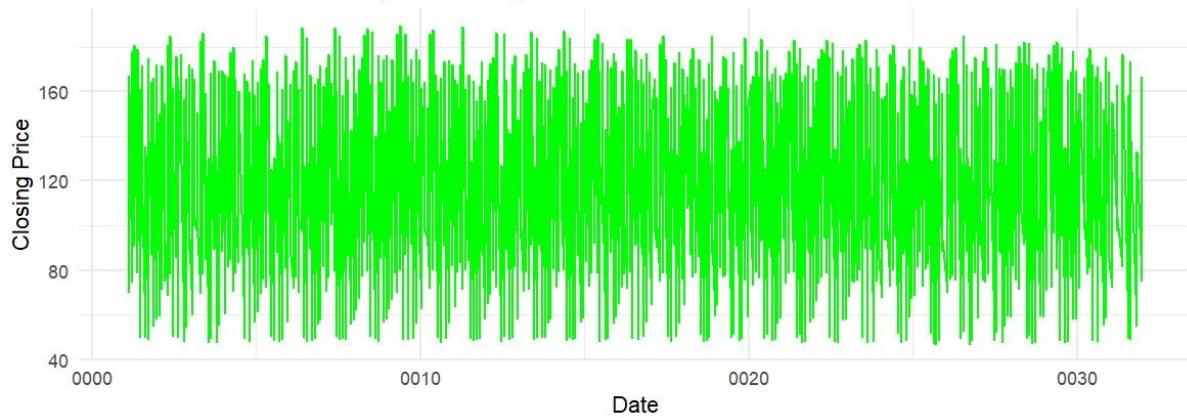
**Output:**



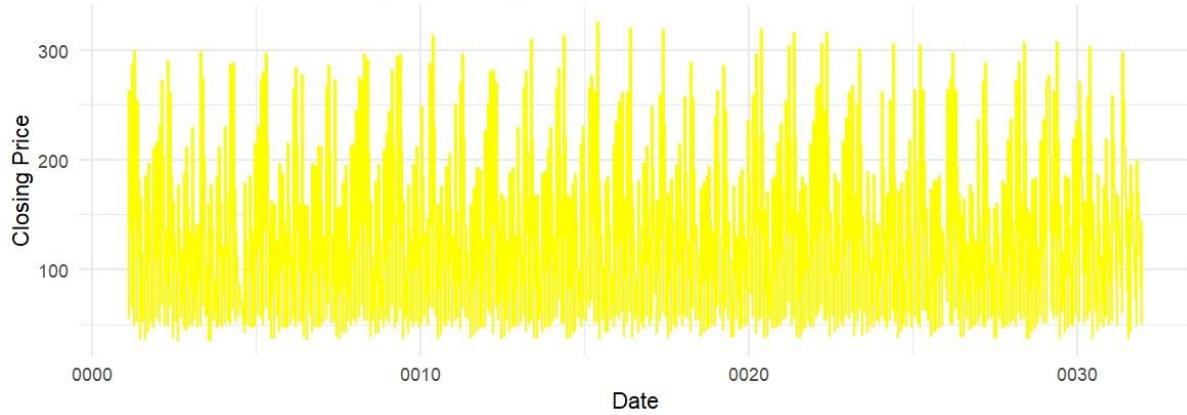
AMD Stock Price Trend (2017–2024)



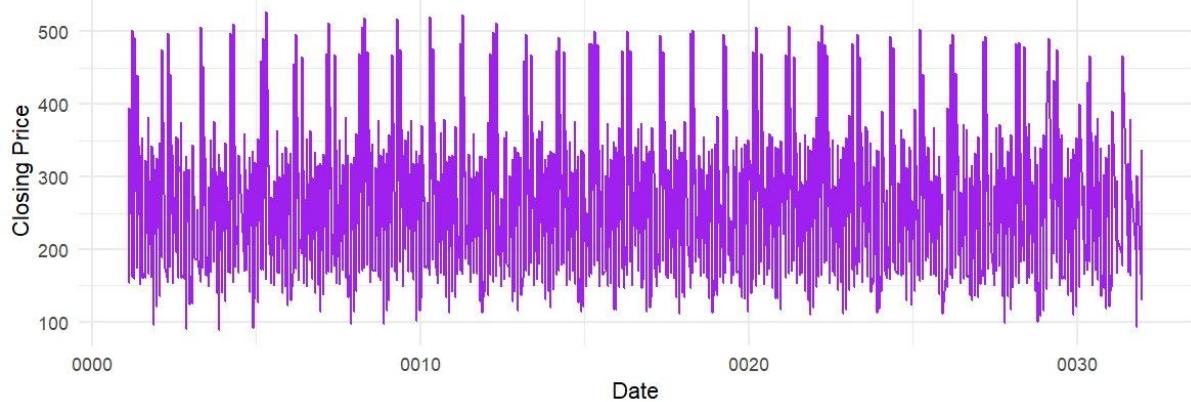
AMZN Stock Price Trend (2017–2024)



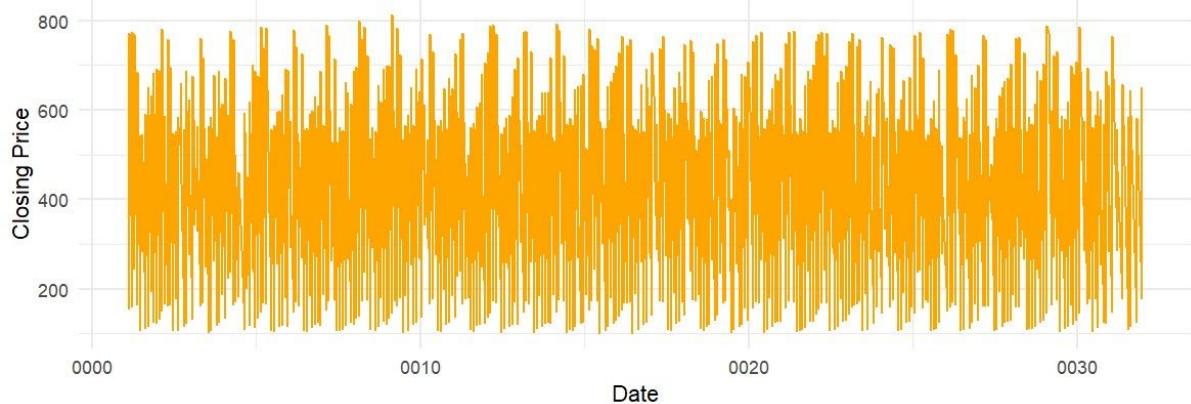
ANET Stock Price Trend (2017–2024)



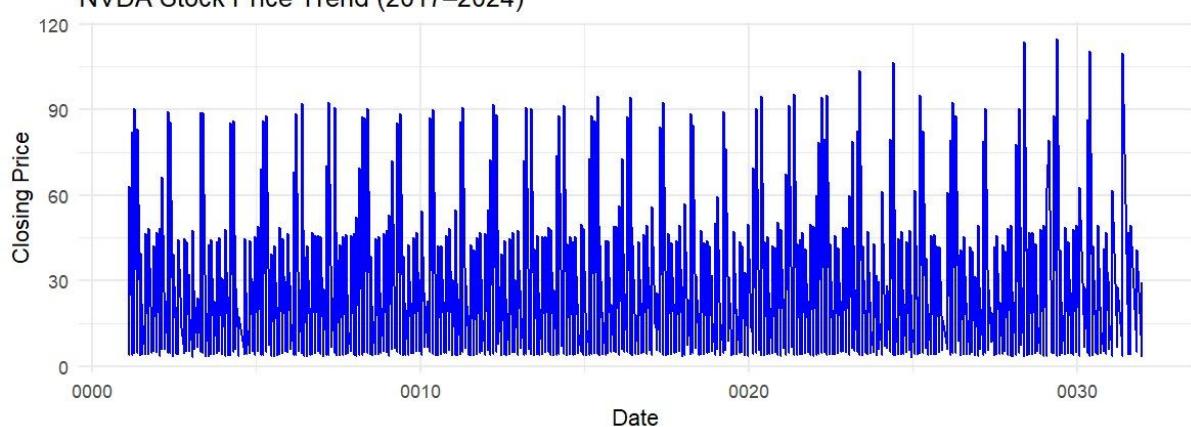
META Stock Price Trend (2017–2024)



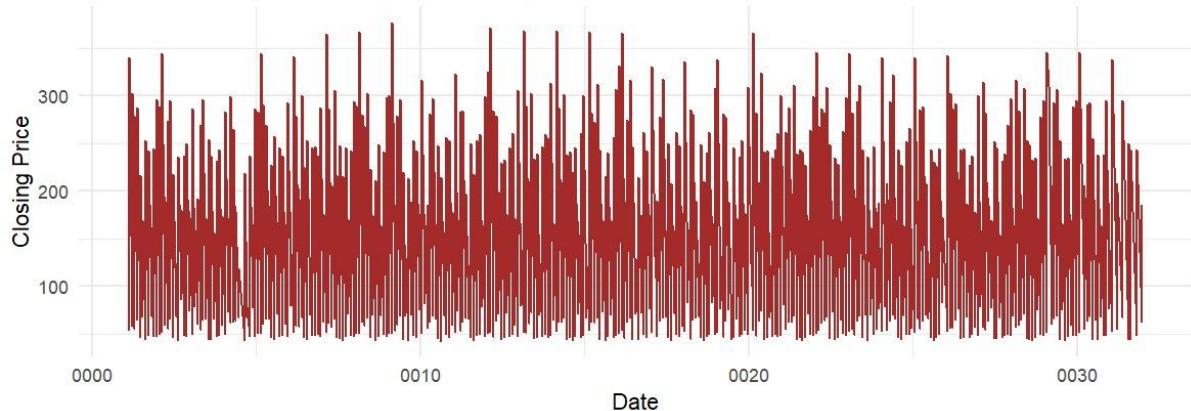
NOW Stock Price Trend (2017–2024)



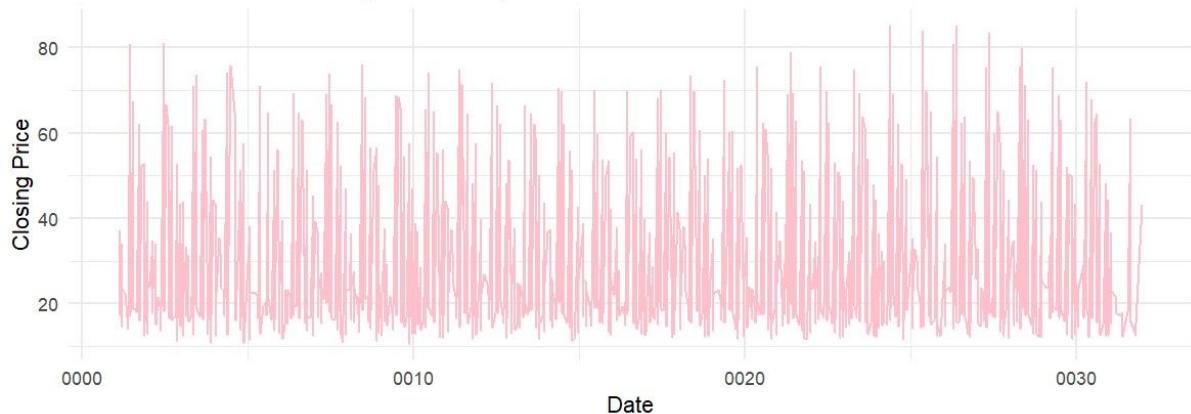
NVDA Stock Price Trend (2017–2024)



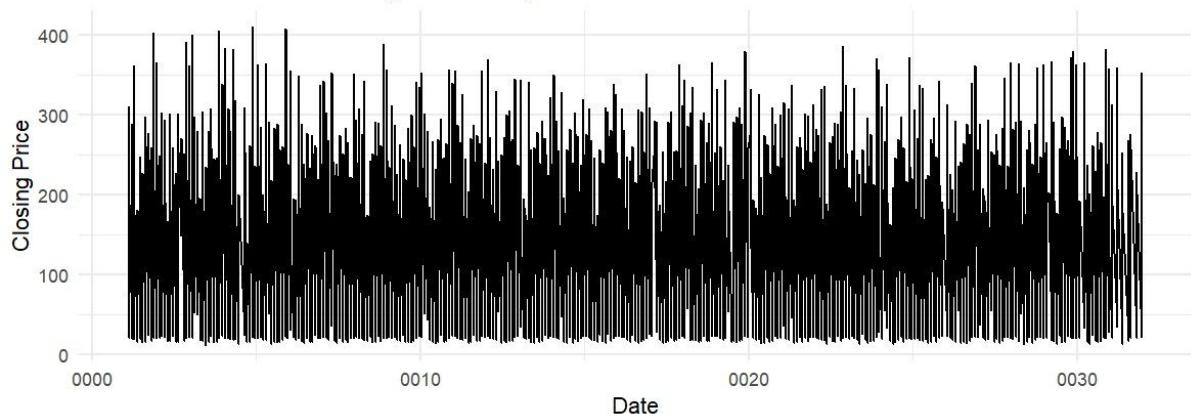
PANW Stock Price Trend (2017–2024)



PATH Stock Price Trend (2017–2024)



TESLA Stock Price Trend (2017–2024)



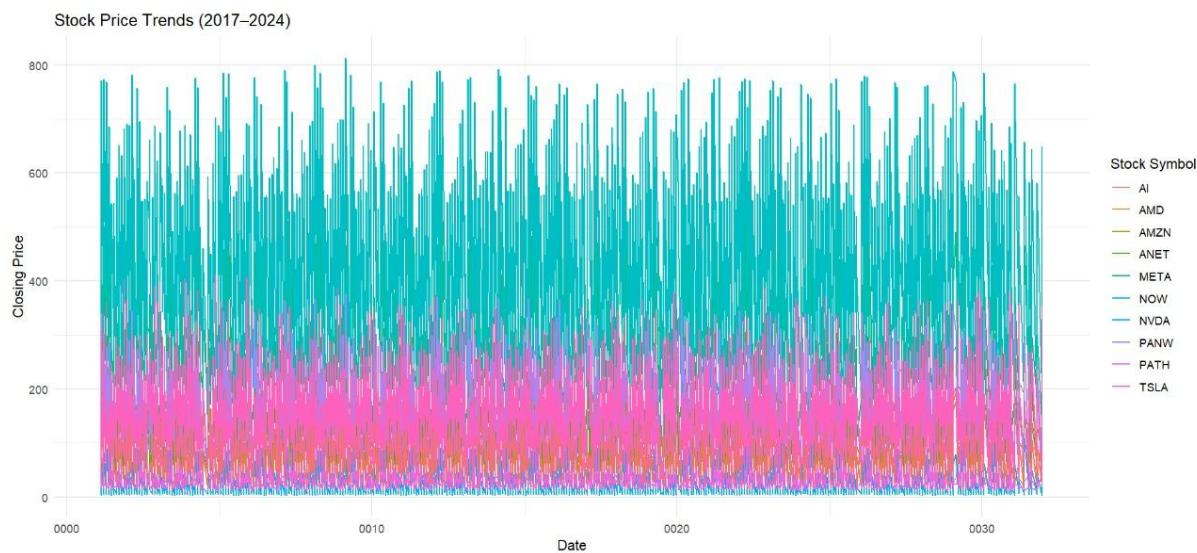
## A. Basic Summary Statistics.

```
87 library(ggplot2)
88
89 ggplot(stock_data, aes(x = date, y = close, color = symbol)) +
90   geom_line() +
91   labs(title = "Stock Price Trends (2017-2024)",
92       x = "Date", y = "Closing Price",
93       color = "Stock Symbol") +
94   theme_minimal()
95
96
97
98 summary_stats <- stock_data %>%
99   group_by(symbol) %>%
100  summarise(
101    Min = min(close, na.rm = TRUE),
102    Max = max(close, na.rm = TRUE),
103    Mean = mean(close, na.rm = TRUE),
104    Median = median(close, na.rm = TRUE),
105    SD = sd(close, na.rm = TRUE),
106    .groups = 'drop'
107  )
108
109 print(summary_stats)
110
111 ggplot(summary_stats, aes(x = reorder(symbol, -Mean), y = Mean, fill = symbol)) +
112   geom_bar(stat = "identity", show.legend = FALSE) +
113   labs(title = "Average Closing Price by Stock (2017-2024)",
114       x = "Stock Symbol", y = "Average Price") +
115   theme_minimal()
```

### Output:

- Consolidated Trends

A consolidated line plot was generated to compare trends across all 10 stocks in one view.



### A. Basic Summary Statistics

```
Descriptive statistics by group
group: AI
  vars   n   mean     sd median trimmed   mad   min   max range skew kurtosis
X1     1 874 37.99 30.33  27.26   31.64 13.28 10.26 177.47 167.21 2.34      5.5
      se
X1 1.03

-----
group: AMD
  vars   n   mean     sd median trimmed   mad   min   max range skew kurtosis
X1     1 1762 69.3 45.86    74.7   65.41 57.49 9.53 211.38 201.85 0.5     -0.48
      se
X1 1.09

-----
group: AMZN
  vars   n   mean     sd median trimmed   mad   min   max range skew kurtosis
X1     1 1762 117.51 39.01 110.89    118 45.08 46.93 189.5 142.57 0.05     -1.19
      se
X1 0.93

-----
group: ANET
  vars   n   mean     sd median trimmed   mad   min   max range skew kurtosis
X1     1 1762 100.79 61.67  73.14   89.24 31.7 35.87 326.03 290.16 1.6      2.06
      se
X1 1.47

-----
group: META
  vars   n   mean     sd median trimmed   mad   min   max range skew kurtosis
X1     1 1762 234.39 87.23 199.8  223.86 60.76 88.91 527.34 438.43 1.15     0.98
      se
X1 2.08
```

```

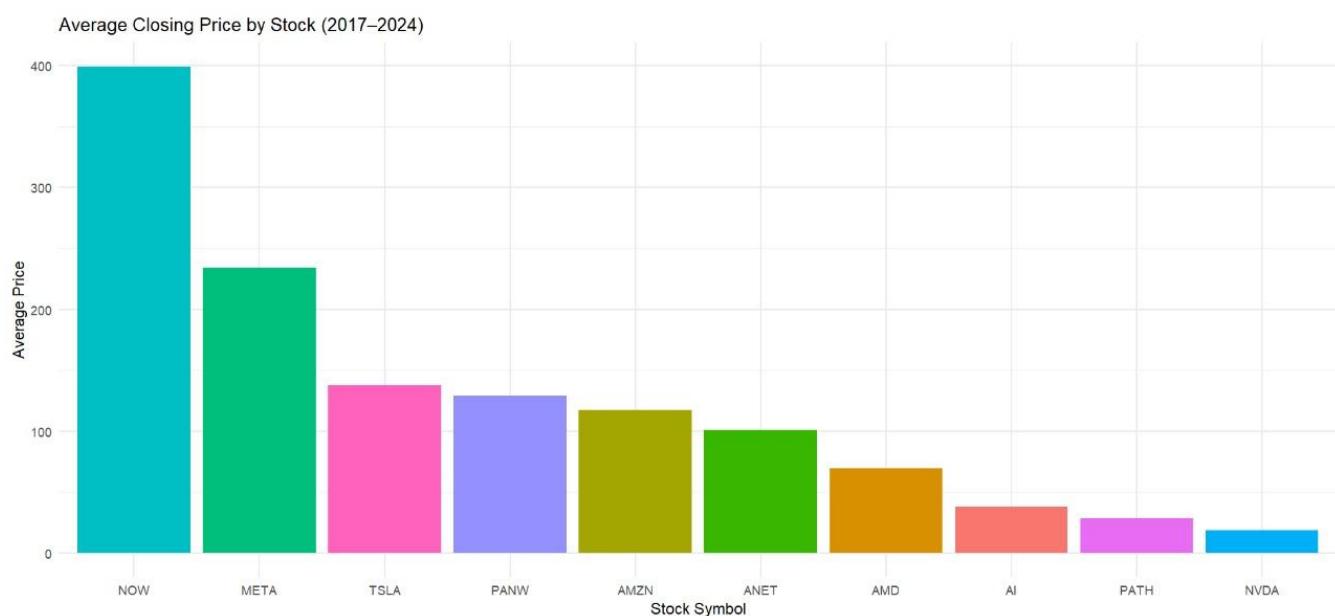
-----
group: NOW
  vars   n   mean      sd median trimmed    mad   min   max range skew
X1     1 1762 399.19 188.84   425.2  393.76 230.78 101.2 812.94 711.74 0.11
  kurtosis se
X1     -1.08 4.5
-----
group: NVDA
  vars   n   mean      sd median trimmed    mad   min   max range skew kurtosis
X1     1 1762 18.65 19.8   12.99   14.47 11.57 3.18 114.82 111.65 2.18      4.95
  se
X1 0.47
-----
group: PANW
  vars   n   mean      sd median trimmed    mad   min   max range skew kurtosis
X1     1 1762 128.92 75.81   98.56  118.93 64.36 42.57 376.9 334.33 0.97      0.11
  se
X1 1.81
-----
group: PATH
  vars   n   mean      sd median trimmed    mad   min   max range skew kurtosis   se
X1     1 784 28.5 18.67   19.41   25.34 6.96 10.55 85.12 74.57 1.3      0.37 0.67
-----
group: TSLA
  vars   n   mean      sd median trimmed    mad   min   max range skew
X1     1 1762 137.98 111.22  144.55  130.03 177.07 11.93 409.97 398.04 0.27
  kurtosis se
X1     -1.33 2.65

```

### 1) Average Closing Price

A bar chart of average prices further supported these insights.

#### Ouput:

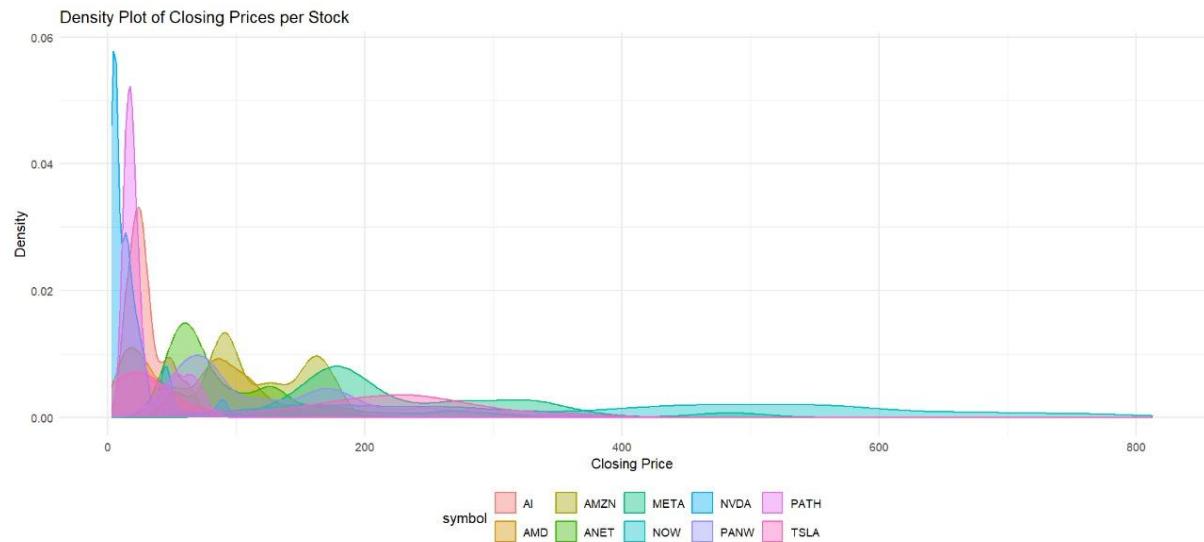


## 2) Density Plot of Closing Prices:

Density plots were created to understand the distribution of closing prices for each stock

```
117 colSums(is.na(stock_data))
118 stock_data <- na.omit(stock_data)
119
120 install.packages("psych") |
121 library(psych)
122
123 describeBy(stock_data$close, group = stock_data$symbol)
124
125 library(ggplot2)
126
127 ggplot(stock_data, aes(x = close, fill = symbol, color = symbol)) +
128   geom_density(alpha = 0.4) +
129   labs(title = "Density Plot of Closing Prices per Stock",
130       x = "Closing Price",
131       y = "Density") +
132   theme_minimal() +
133   theme(legend.position = "bottom")
```

**Output:**



## 1) Correlation heatmap for Stock Closing prices:

```
233 library(ggplot2)
234 library(reshape2)
235 library(viridis)
236 cor_matrix <- cor(wide_data_clean[-1], use = "complete.obs")
237 melted_corr <- melt(cor_matrix)
238 ggplot(melted_corr, aes(x = Var1, y = Var2, fill = value)) +
239   geom_tile(color = "white") +
240   scale_fill_viridis(name = "Correlation", limits = c(-1, 1)) +
241   labs(
242     title = "Correlation Heatmap of Stock Closing Prices",
243     x = "Stock Symbol",
244     y = "Stock Symbol"
245   ) +
246   theme_minimal(base_size = 13) +
247   theme(axis.text.x = element_text(angle = 45, vjust = 1, hjust = 1))
```

## Output:



## 2) Time Series Decomposition

To analyze the underlying structure of stock price movements over time, STL (Seasonal-Trend decomposition using Loess) was applied to each stock's closing price data. This approach helped separate each time series into three distinct components:

- Trend: Long-term progression of the series
- Seasonal: Repeating short-term patterns (e.g., yearly cycles)
- Remainder: Residual variation (noise or unexplained fluctuations)

We used the `stl()` function from the `forecast` package in R with a periodic seasonal window and daily frequency (frequency = 365) assuming annual seasonality.

```

249 library(forecast)
250 library(ggplot2)
251 nvda_data <- raw_data %>%
252   filter(symbol == "NVDA") %>%
253   mutate(date = as.Date(date, format = "%d-%m-%Y")) %>%
254   arrange(date)
255 nvda_ts <- ts(nvda_data$close, frequency = 365, start = c(2017, 1))
256 nvda_decomp <- stl(nvda_ts, s.window = "periodic")
257 autoplot(nvda_decomp) +
258   labs(title = "Time Series Decomposition: NVDA", x = "Time")
259
260 ai_data <- raw_data %>%
261   filter(symbol == "AI") %>%
262   mutate(date = as.Date(date, format = "%d-%m-%Y")) %>%
263   arrange(date)
264 ai_ts <- ts(ai_data$close, frequency = 365, start = c(2017, 1))
265 ai_decomp <- stl(ai_ts, s.window = "periodic")
266 autoplot(ai_decomp) +
267   labs(title = "Time Series Decomposition: AI", x = "Time")
268
269
270
271 amd_data <- raw_data %>%
272   filter(symbol == "AMD") %>%
273   mutate(date = as.Date(date, format = "%d-%m-%Y")) %>%
274   arrange(date)
275 amd_ts <- ts(amd_data$close, frequency = 365, start = c(2017, 1))
276 amd_decomp <- stl(amd_ts, s.window = "periodic")
277 autoplot(amd_decomp) +
278   labs(title = "Time Series Decomposition: AMD", x = "Time")

```

```

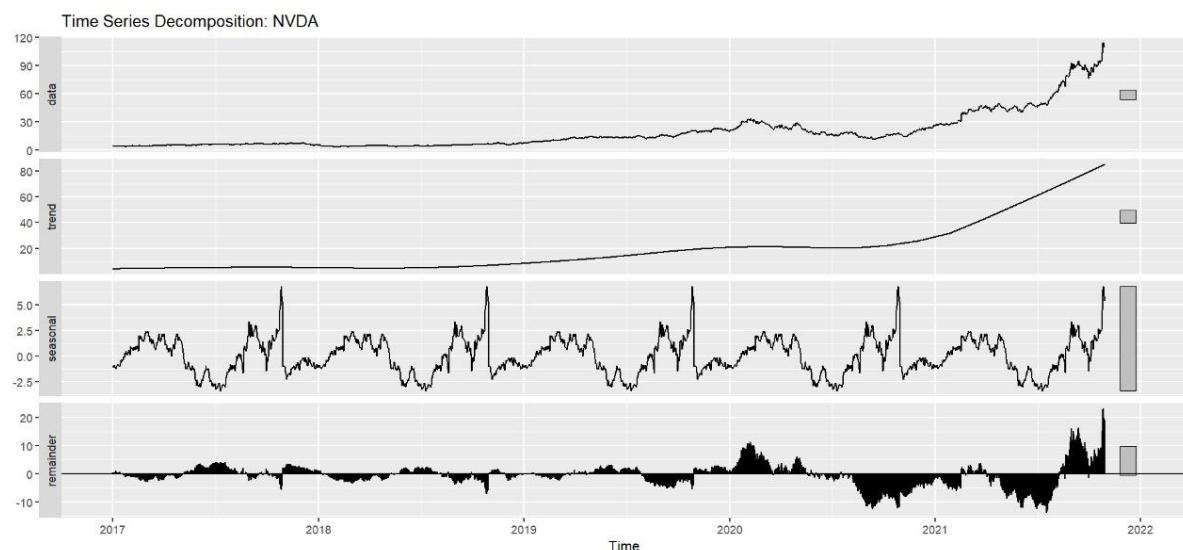
298 path_data <- raw_data %>%
299   filter(symbol == "PATH") %>%
300   mutate(date = as.Date(date, format = "%d-%m-%Y")) %>%
301   arrange(date)
302 path_ts <- ts(path_data$close, frequency = 365, start = c(2017, 1))
303 path_decomp <- stl(path_ts, s.window = "periodic")
304 autoplot(path_decomp) +
305   labs(title = "Time Series Decomposition: PATH", x = "Time")
306
307 anet_data <- raw_data %>%
308   filter(symbol == "ANET") %>%
309   mutate(date = as.Date(date, format = "%d-%m-%Y")) %>%
310   arrange(date)
311
312 anet_ts <- ts(anet_data$close, frequency = 365, start = c(2017, 1))
313 anet_decomp <- stl(anet_ts, s.window = "periodic")
314 autoplot(anet_decomp) +
315   labs(title = "Time Series Decomposition: ANET", x = "Time")
316
317 meta_data <- raw_data %>%
318   filter(symbol == "META") %>%
319   mutate(date = as.Date(date, format = "%d-%m-%Y")) %>%
320   arrange(date)
321
322 meta_ts <- ts(meta_data$close, frequency = 365, start = c(2017, 1))
323 meta_decomp <- stl(meta_ts, s.window = "periodic")
324 autoplot(meta_decomp) +
325   labs(title = "Time Series Decomposition: META", x = "Time")

```

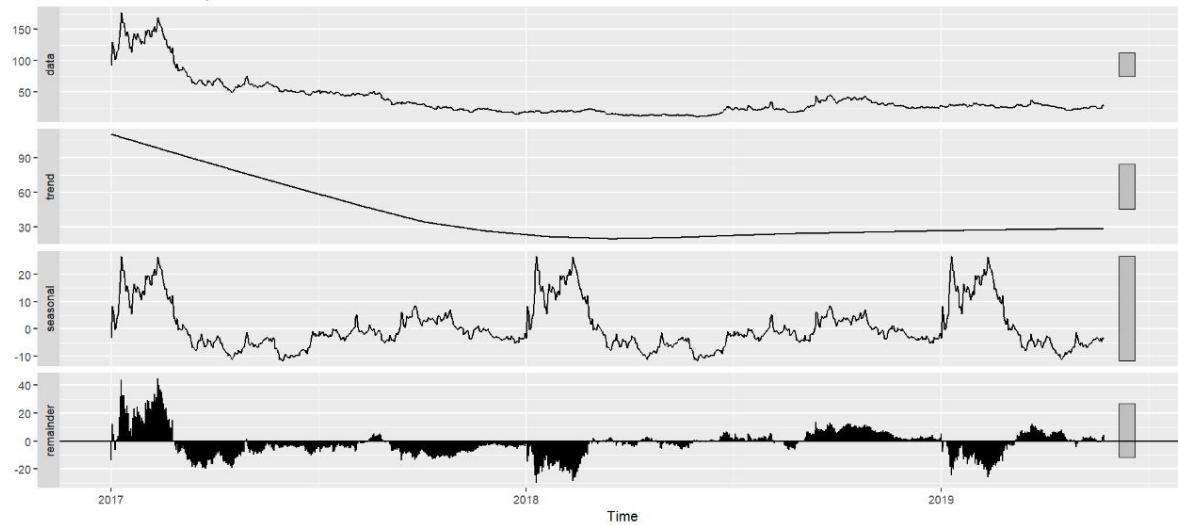
```

327 amzn_data <- raw_data %>%
328   filter(symbol == "AMZN") %>%
329   mutate(date = as.Date(date, format = "%d-%m-%Y")) %>%
330   arrange(date)
331
332 amzn_ts <- ts(amzn_data$close, frequency = 365, start = c(2017, 1))
333 amzn_decomp <- stl(amzn_ts, s.window = "periodic")
334 autoplot(amzn_decomp) +
335   labs(title = "Time Series Decomposition: AMZN", x = "Time")
336
337 now_data <- raw_data %>%
338   filter(symbol == "NOW") %>%
339   mutate(date = as.Date(date, format = "%d-%m-%Y")) %>%
340   arrange(date)
341
342 now_ts <- ts(now_data$close, frequency = 365, start = c(2017, 1))
343 now_decomp <- stl(now_ts, s.window = "periodic")
344 autoplot(now_decomp) +
345   labs(title = "Time Series Decomposition: NOW", x = "Time")

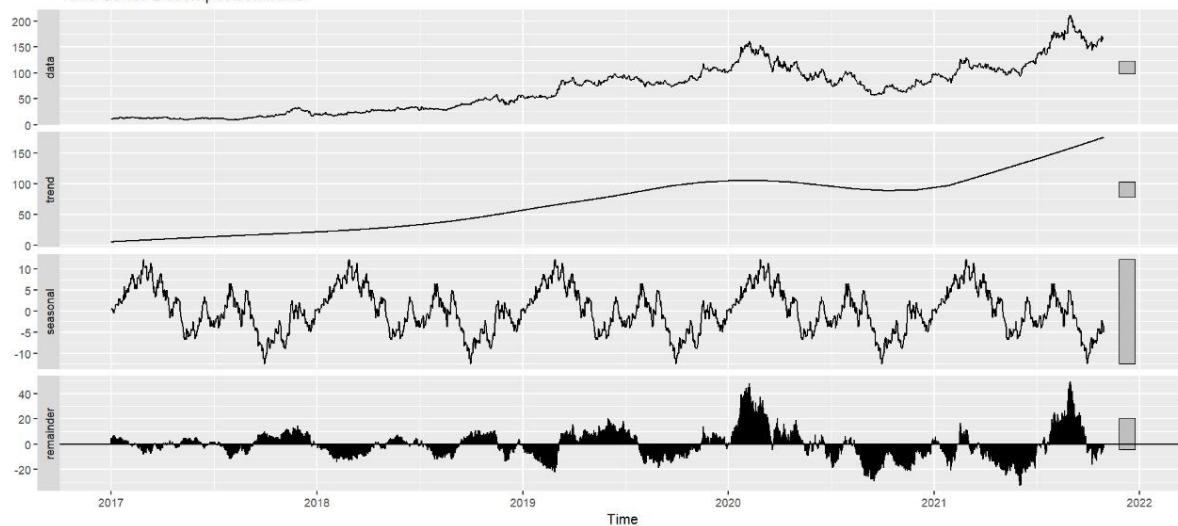
```



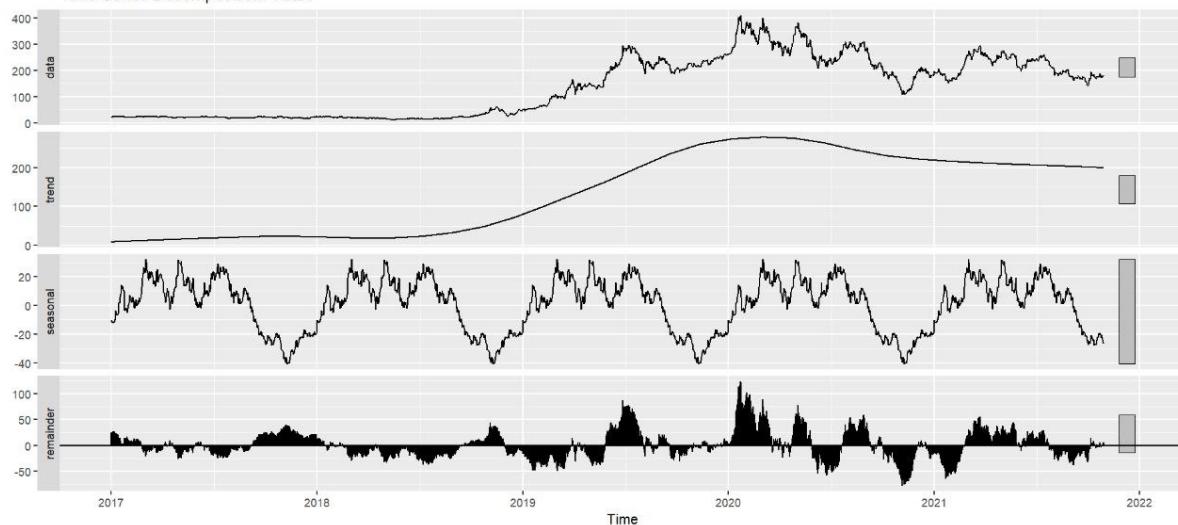
Time Series Decomposition: AI



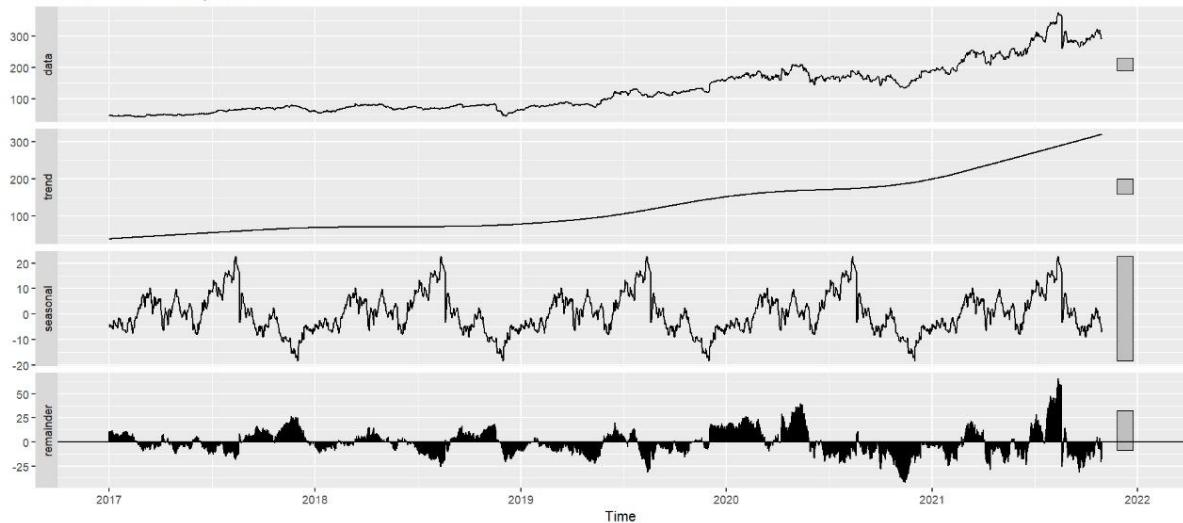
Time Series Decomposition: AMD



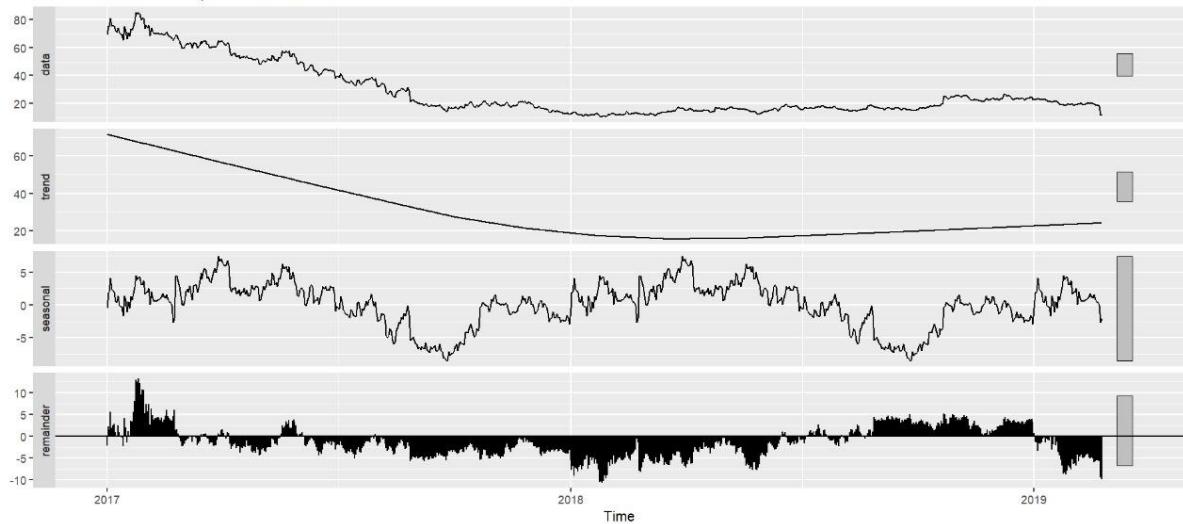
Time Series Decomposition: TSLA



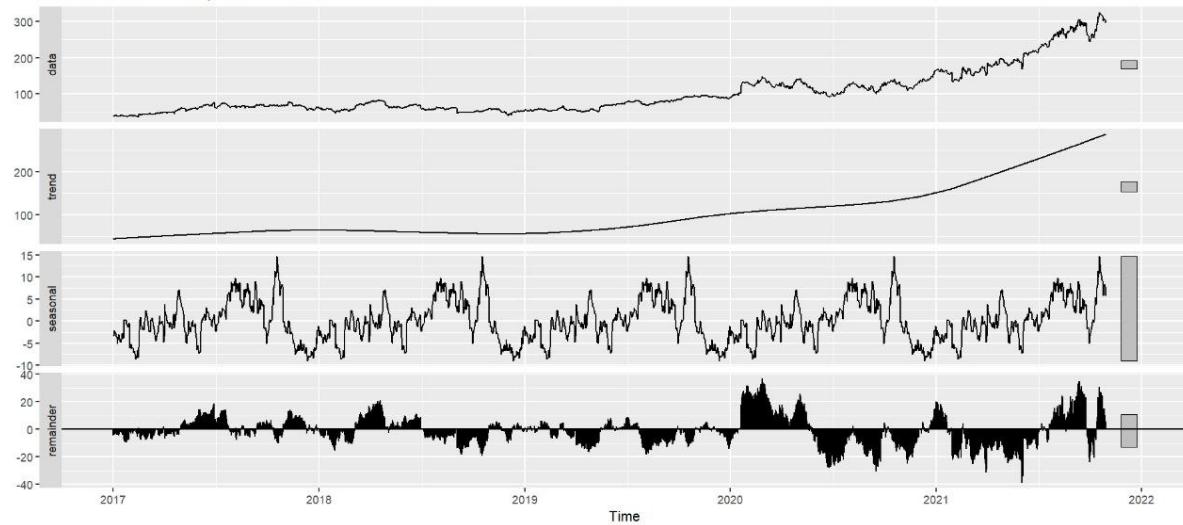
Time Series Decomposition: PANW



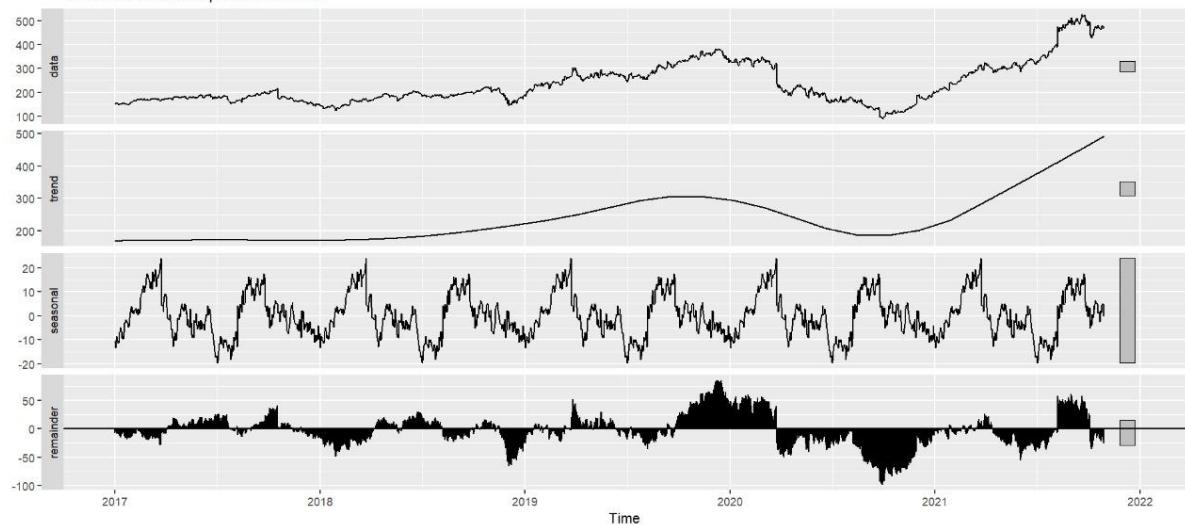
Time Series Decomposition: PATH



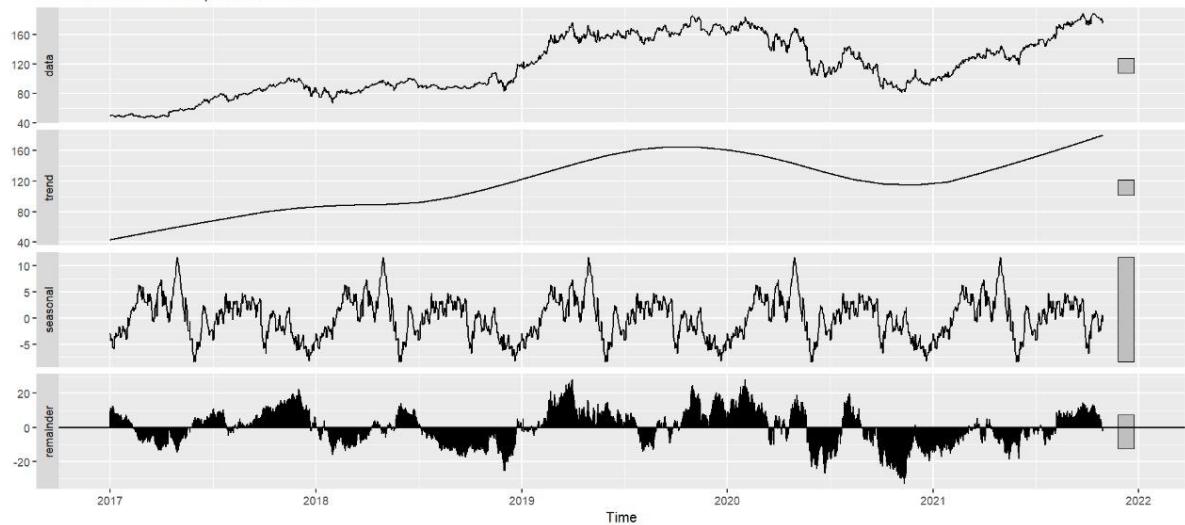
Time Series Decomposition: ANET



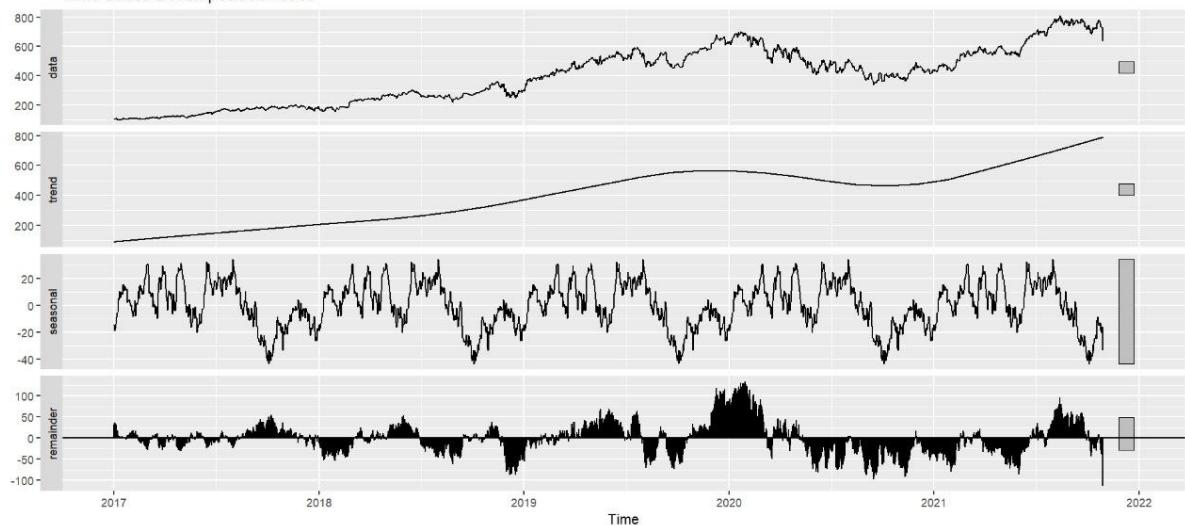
Time Series Decomposition: META



Time Series Decomposition: AMZN



Time Series Decomposition: NOW



## B. Forecasting

### 1) ARIMA Forecasting:

To model and predict short-term future stock prices, we applied the ARIMA (AutoRegressive Integrated Moving Average) technique to each of the ten AI-related stocks from 2017 to 2024.

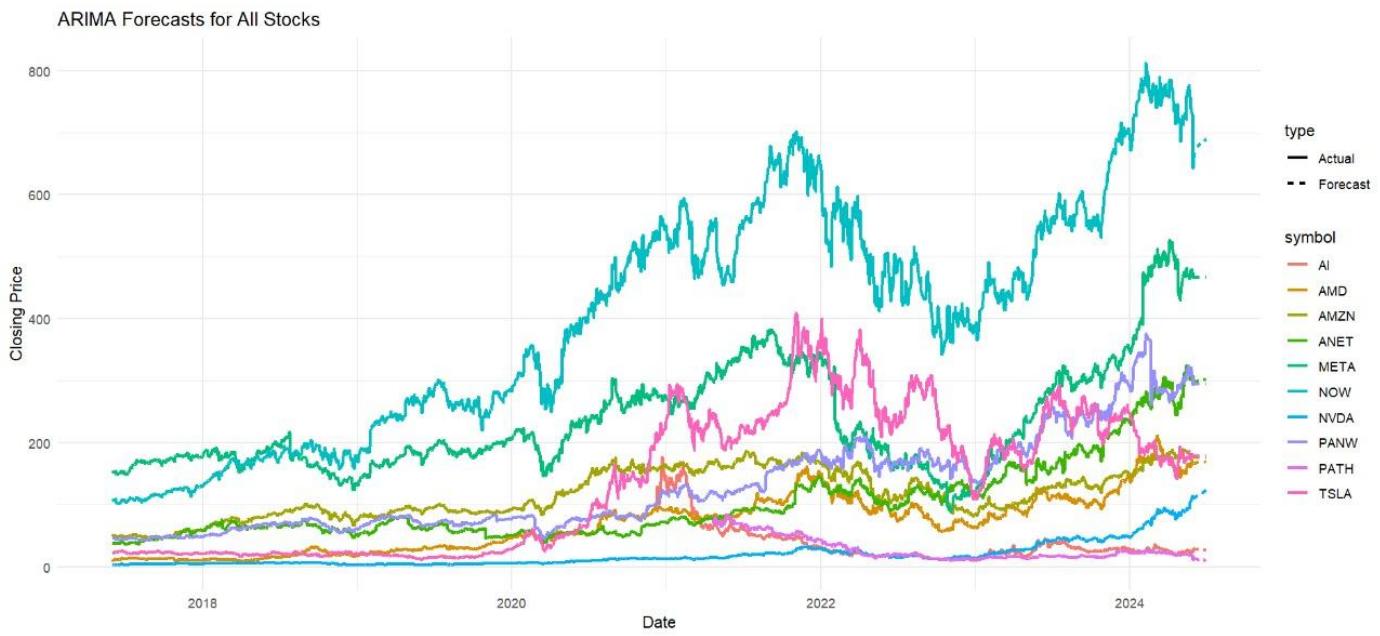
Using the `auto.arima()` function from the `forecast` package in R, we automatically selected the best-fitting ARIMA model for each stock's time series. Each model was then used to generate a 30-day forecast of closing prices beyond the last available date in the dataset.

```
355 library(forecast)
356 library(ggplot2)
357 library(dplyr)
358 library(tidyr)
359
360 raw_data <- raw_data %>%
361   mutate(date = as.Date(date, format = "%d-%m-%Y"))
362 stock_list <- unique(raw_data$symbol)
363
364 forecast_list <- list()
365 for (stock in stock_list) {
366   stock_data <- raw_data %>%
367     filter(symbol == stock) %>%
368     arrange(date)
369   ts_data <- ts(stock_data$close, frequency = 365)
370   model <- auto.arima(ts_data)
371   forecast_result <- forecast(model, h = 30)
372   last_date <- max(stock_data$date)
373   forecast_dates <- seq.Date(from = last_date + 1, by = "day", length.out = 30)
374   forecast_df <- data.frame(
375     date = forecast_dates,
376     close = as.numeric(forecast_result$mean),
377     symbol = stock,
378     type = "Forecast"
379   )
380   actual_df <- stock_data %>%
381     select(date, close) %>%
382     mutate(symbol = stock, type = "Actual")
383
384   forecast_list[[stock]] <- rbind(actual_df, forecast_df)
385 }
```

```
387 all_forecast_df <- do.call(rbind, forecast_list)
388 ggplot(all_forecast_df, aes(x = date, y = close, color = symbol, linetype = type)) +
389   geom_line(size = 1) +
390   labs(title = "ARIMA Forecasts for All Stocks", x = "Date", y = "Closing Price") +
391   theme_minimal()
```

## Output:



## 2) LSTM Forecasting using Torch in R:

To capture complex temporal dependencies and nonlinear trends in stock prices, we implemented a Long Short-Term Memory (LSTM) model using the torch package in R. LSTMs are a class of recurrent neural networks (RNNs) especially well-suited for time series forecasting due to their ability to retain long-term information.

- **Data Preparation:**

For each stock symbol, data was filtered and sorted by date.  
A 5-day moving average was computed to smooth short-term fluctuations.  
The resulting series was normalized to [0, 1] for stability in training.

- **Sequence Construction:**

For each stock, sliding sequences of length 50 were generated (seq\_length = 50).  
Each input sequence predicted the next value in the series (supervised learning setup).

- **Train-Test Split:**

80% of the data was used for training; 20% for testing the model's generalization.

- **Model Architecture:**

The LSTM model consisted of:

- 2 stacked LSTM layers with 64 hidden units.
- Dropout for regularization ( $p = 0.2$ ).
- A fully connected output layer to predict the next value.
- The loss function used was Mean Squared Error (MSE), optimized using Adam.

- **Training:**

Each model was trained for 50 epochs with a learning rate of 0.001.  
Training loss was monitored every 10 epochs for convergence diagnostics.

- **Evaluation:**

Predictions were de-normalized to obtain values in the original price scale.  
Performance metrics computed for each stock included:

- MAE (Mean Absolute Error)
- RMSE (Root Mean Square Error)
- MAPE (Mean Absolute Percentage Error)
- SMAPE (Symmetric Mean Absolute Percentage Error)

Visual comparisons of actual vs. predicted prices were plotted.

```

1099 library(torch)
1100 library(dplyr)
1101 library(ggplot2)
1102 library(zoo)
1103
1104 # Parameters
1105 seq_length <- 50
1106 epochs <- 50
1107 learning_rate <- 0.001
1108
1109 # Loop through each unique stock
1110 unique_stocks <- unique(raw_data$symbol)
1111
1112 for (stock in unique_stocks) {
1113   cat("Training for stock:", stock, "\n")
1114
1115   stock_data <- raw_data %>% filter(symbol == stock)
1116   if (nrow(stock_data) < 200) {
1117     cat("Skipping", stock, "- not enough data\n")
1118     next
1119   }
1120
1121 # Smooth close prices using rolling average
1122 series <- zoo::rollmean(stock_data$close, k = 5, fill = NA, align = "right")
1123 series <- na.omit(series)
1124
1125 # Normalize
1126 min_val <- min(series)
1127 max_val <- max(series)
1128 series_scaled <- (series - min_val) / (max_val - min_val)
1129 series_tensor <- torch_tensor(as.numeric(series_scaled), dtype = torch_float())
1130

```

```

1131 # Create sequences
1132 x <- list()
1133 y <- list()
1134 for (i in 1:(length(series_tensor) - seq_length)) {
1135   x[[i]] <- series_tensor[i:(i + seq_length - 1)]
1136   y[[i]] <- series_tensor[i + seq_length]
1137 }
1138
1139 x <- torch_stack(x) %>% torch_unsqueeze(3)
1140 y <- torch_stack(y)
1141
1142 # Train-test split (80-20)
1143 n <- x$csize(1)
1144 train_size <- floor(0.8 * n)
1145 x_train <- x[1:train_size,,]
1146 y_train <- y[1:train_size]
1147 x_test <- x[(train_size+1):n,,]
1148 y_test <- y[(train_size+1):n]
1149
1150 # Model
1151 model <- nn_module(
1152   "LSTMModel",
1153   initialize = function() {
1154     self$lstm <- nn_lstm(input_size = 1, hidden_size = 64, num_layers = 2, batch_first = TRUE)
1155     self$dropout <- nn_dropout(p = 0.2)
1156     self$fc <- nn_linear(64, 1)
1157   },
1158   forward = function(x) {
1159     out <- self$lstm(x)
1160     h_n <- out[[1]][, dim(out[[1]])[2], ]
1161     h_n <- self$dropout(h_n)
1162     self$fc(h_n)$squeeze(2)
1163   }
1164 )

```

```

1165 net <- model()
1166 net$train()
1167 optimizer <- optim_adam(net$parameters, lr = learning_rate)
1168 loss_fn <- nn_mse_loss()
1169
1170
1171 for (epoch in 1:epochs) {
1172   optimizer$zero_grad()
1173   output <- net(x_train)
1174   loss <- loss_fn(output, y_train)
1175   loss$backward()
1176   optimizer$step()
1177   if (epoch %% 10 == 0) {
1178     cat("Epoch:", epoch, "Loss:", loss$item(), "\n")
1179   }
1180 }
1181
1182 # Evaluation
1183 net$eval()
1184 predicted <- with_no_grad({
1185   net(x_test)
1186 })
1187
1188 predicted_vals <- as.numeric(predicted) * (max_val - min_val) + min_val
1189 actual_vals <- as.numeric(y_test) * (max_val - min_val) + min_val
1190
1191 # Plot
1192 df <- data.frame(
1193   Time = 1:length(predicted_vals),
1194   Predicted = predicted_vals,
1195   Actual = actual_vals
1196 )

```

```

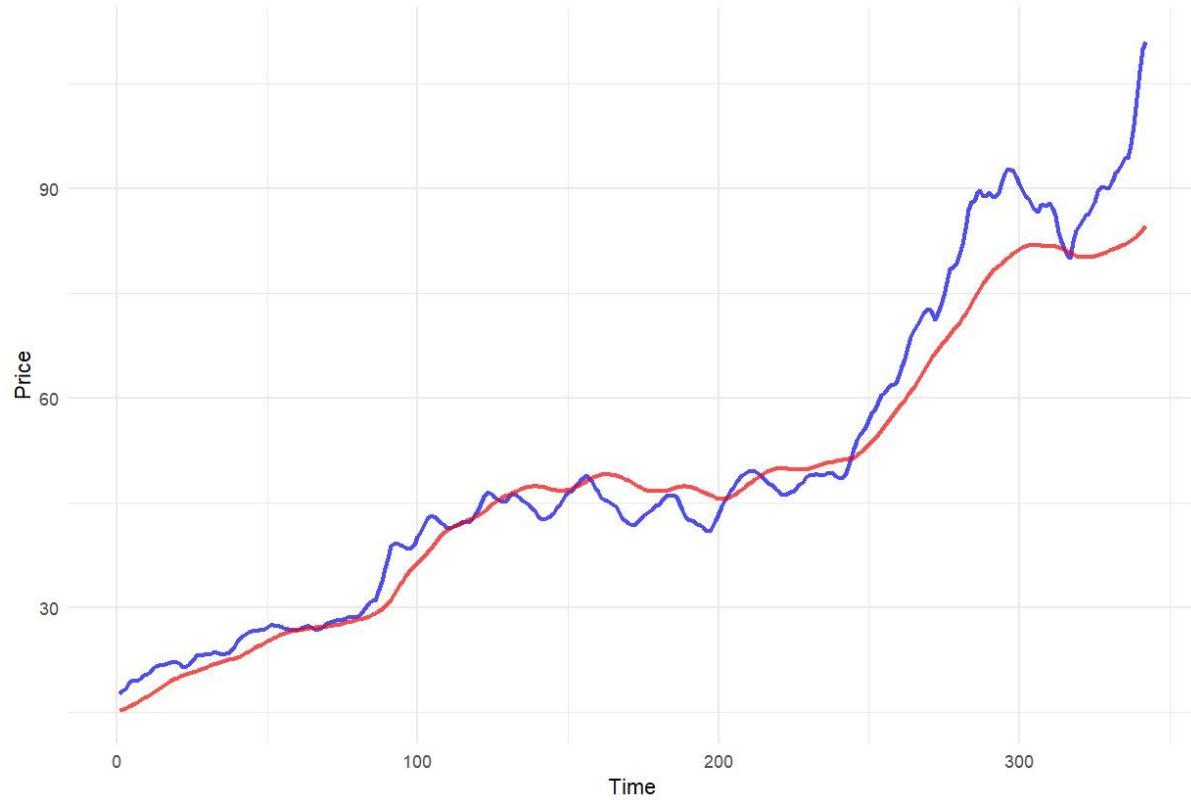
1198 plot <- ggplot(df, aes(x = Time)) +
1199   geom_line(aes(y = Actual), color = "blue", size = 1.1, alpha = 0.7) +
1200   geom_line(aes(y = Predicted), color = "red", size = 1.1, alpha = 0.7) +
1201   labs(title = paste("Prediction vs Actual for", stock), y = "Price", x = "Time") +
1202   theme_minimal()
1203
1204 print(plot)
1205
1206 # Error metrics
1207 mae <- mean(abs(actual_vals - predicted_vals))
1208 rmse <- sqrt(mean((actual_vals - predicted_vals)^2))
1209 mape <- mean(abs((actual_vals - predicted_vals) / actual_vals)) * 100
1210 smape <- mean(200 * abs(predicted_vals - actual_vals) / (abs(actual_vals) + abs(predicted_vals))) * 100
1211
1212 cat("Stock:", stock, "\n")
1213 cat("MAE:", round(mae, 5), "\n")
1214 cat("RMSE:", round(rmse, 5), "\n")
1215 cat("MAPE:", round(mape, 2), "%\n")
1216 cat("SMAPE:", round(smape, 2), "%\n\n")
1217 }

```

**Output:**

The predicted outputs of few selected stocks:

Prediction vs Actual for NVDA



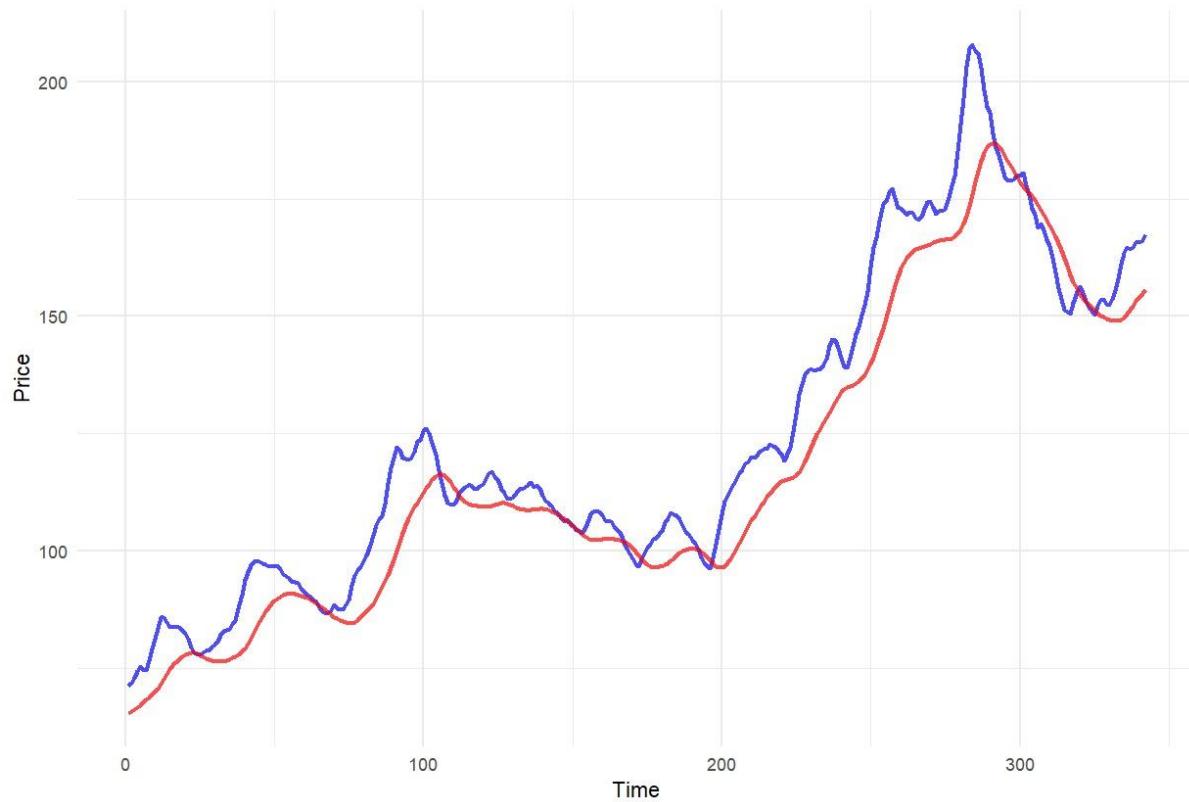
Prediction vs Actual for AMZN



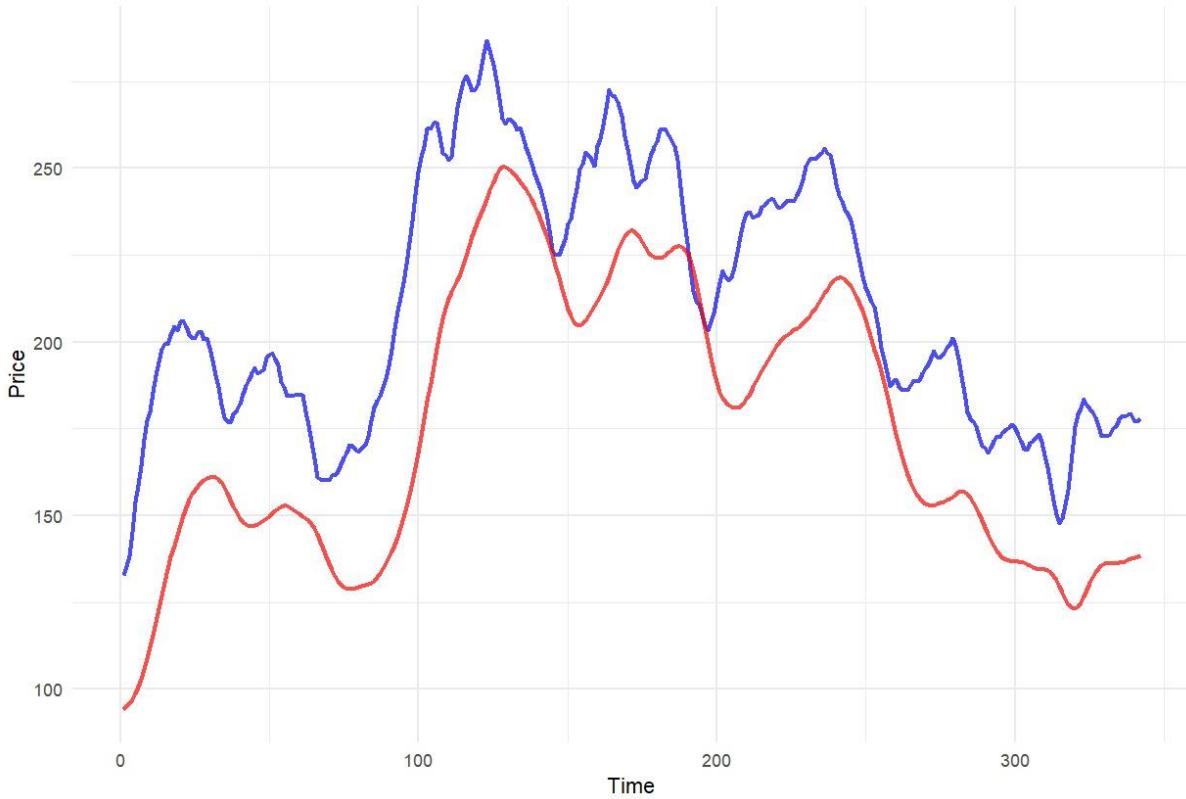
Prediction vs Actual for AI



Prediction vs Actual for AMD



Prediction vs Actual for TSLA



### C. Volatility Analysis using 30-Day Rolling Standard Deviation

To analyze the volatility behavior of each stock over time, we computed the 30-day rolling standard deviation of closing prices using the zoo package in R. Rolling standard deviation serves as a proxy for volatility, highlighting periods of price stability or turbulence.

- **Methodology:**

The dataset was grouped by stock symbol and sorted by date to ensure chronological consistency.

For each stock, a 30-day window was applied to compute the standard deviation of closing prices using rollapply():

$$\text{Volatility}_t = \text{SD}(P_{t-29}, \dots, P_t)$$

where  $P_t$  is the closing price on day  $t$

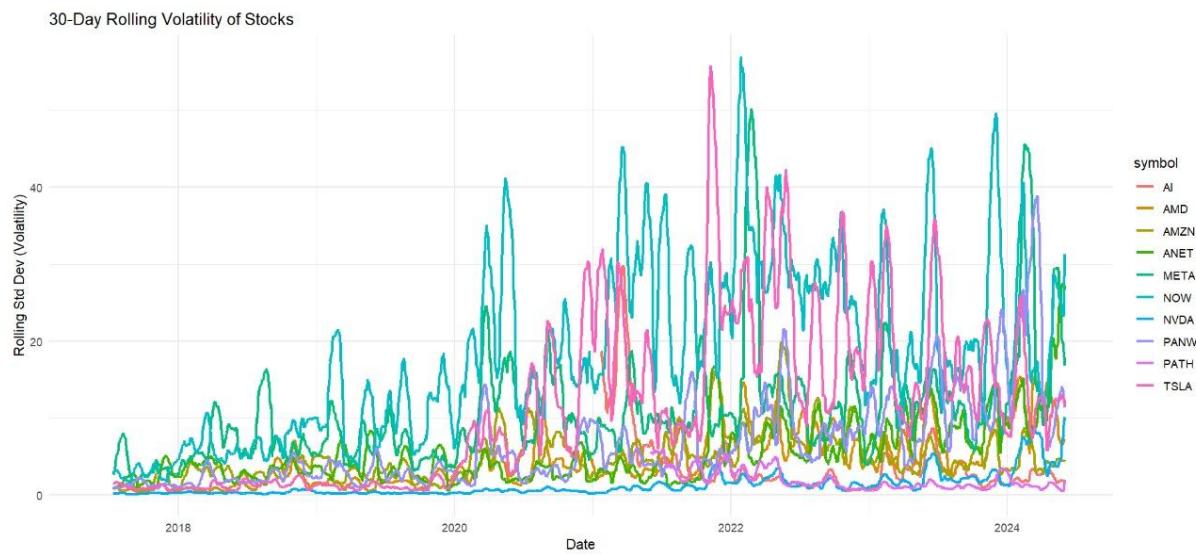
The result (rolling\_sd) captures local variations in price levels, reflecting short-term volatility trends.

```

385 library(zoo)
386
387 # Calculate 30-day rolling standard deviation for each stock
388 rolling_volatility <- raw_data %>%
389   arrange(symbol, date) %>%
390   group_by(symbol) %>%
391   mutate(rolling_sd = rollapply(close, width = 30, FUN = sd, fill = NA, align = "right"))
392
393 # Plot all stocks' volatility together
394 ggplot(rolling_volatility, aes(x = date, y = rolling_sd, color = symbol)) +
395   geom_line(size = 1) +
396   labs(title = "30-Day Rolling Volatility of Stocks",
397       x = "Date",
398       y = "Rolling Std Dev (Volatility)") +
399   theme_minimal()

```

## Output:



## VI. RESULT AND INFERENCE

### A. Exploratory Data Analysis (EDA)

#### 1. Trend Insights:

- Stocks like NVDA, TSLA, and AMD showed **strong upward trends**, especially after 2020, reflecting the AI boom.
- AI, PATH, and PANW exhibited more **volatile behavior**, with irregular growth patterns.
- META, AMZN, and NOW remained relatively **stable**, though with smaller fluctuations.

#### 2. Summary Statistics:

- NVDA had the **highest average closing price**, indicating its leading market performance in the AI sector.
- TSLA and AMZN showed **high standard deviation**, suggesting greater price fluctuation (risk).
- PATH had one of the **lowest means**, showing it's either newer or under lower market confidence.

#### 3. Average Closing Price Analysis (2017–2024)

- To identify which AI-related stocks had the strongest market value over time, the **average (mean) closing price** was calculated for each stock. This metric provides a high-level view of a company's valuation trajectory across the 7-year period.

Results (Summary Table):

Stock	Min	Max	Mean	Median	SD	Range
AI	10.26	177.47	37.99	27.26	30.33	167.21
AMD	9.53	211.38	69.30	74.70	45.86	201.85

Stock	Min	Max	Mean	Median	SD	Range
AMZN	46.93	189.50	117.51	110.89	39.01	142.57
ANET	35.87	326.03	100.79	73.14	61.67	290.16
META	88.91	527.34	234.39	199.80	87.23	438.43
NOW	101.20	812.94	399.19	425.20	188.84	711.74
NVDA	3.18	114.82	18.65	12.99	19.80	111.65
PANW	42.57	376.90	128.92	98.56	75.81	334.33
PATH	0.37	28.50	18.67	19.41	0.67	28.13
TSLA	11.93	409.97	137.98	144.55	111.22	398.04

Updated Inference:

**NOW** has the **highest average price** (Mean = 399.19) among all stocks, with a very wide range (711.74), indicating significant historical price movement.

**TSLA** also shows **high volatility** with SD = 111.22 and a wide range (398.04), suggesting strong price fluctuations.

**NVDA** and **PATH** have **relatively lower average prices**, with NVDA showing high skew (2.18) suggesting a few very high price values.

**PANW** has a **moderate mean** (128.92) but a **large range** (334.33), highlighting notable volatility.

**PATH** has the **lowest range** and SD, indicating it remained relatively stable in price.

#### 4. Density Plot

*Inference:*

*Stocks like ANET, NOW, and META had tight distributions, indicating more predictable behavior.*

*TSLA and NVDA had long-tailed distributions, reflecting rapid price acceleration phases.*

#### 5. Correlation Heatmap:

*Inference:*

- Strong positive correlations were observed between **AMZN, META, and NVDA**, suggesting similar market behavior.
- Lower or no correlation seen between **TSLA** and others due to its unique price dynamics.

#### 6. Time Series Decomposition

##### Results

For each stock, the closing price time series was decomposed into the following components:

- Trend:** Reflects the general direction of the stock price over time (upward, downward, or flat).
- Seasonality:** Recurring patterns or cycles (e.g., quarterly dips, year-end rallies).
- Residual:** Irregular or random fluctuations not explained by trend or seasonality.

Here's a general overview of the observations based on decomposition

Stock	Trend Behavior	Seasonality Strength	Residual Noise
AI	Gradual upward trend post-2023	Mild	High (volatile early phase)
AMD	Strong upward trend until 2022, then flat	Moderate	Low to moderate
AMZN	Long-term upward trend	Moderate	Low noise
ANET	Upward trend with fluctuation	Weak	Noticeable spikes
META	Volatile but rising trend	Weak	High
NOW	Steady and strong trend	Mild	Low
NVDA	Highly volatile trend	Minimal	Strong residuals
PANW	Seasonal growth pattern	Strong	Medium noise
PATH	Flat to rising in late years	Weak	High
TSLA	Highly nonlinear trend	Weak to moderate	Significant noise

#### Inference from Decomposition

**Long-Term Investors** should focus on stocks with a strong and steady **trend component** (e.g., NOW, AMD, AMZN).

**Seasonality-sensitive stocks** like PANW could benefit from **seasonal trading strategies**—especially if price patterns repeat predictably.

Stocks like NVDA, META, and TSLA have **high residuals**, indicating that forecasting models need to handle more randomness or external influence.

Stocks with **weak seasonality** (e.g., PATH, AI) suggest that calendar-based effects are less relevant.

#### B. Forecasting

##### 1. ARIMA Forecasting

Forecast Behavior:

- **ARIMA models** effectively projected short-term price movements for most stocks.
- Forecast curves for stocks like META and ANET closely followed recent trends, reflecting stable price dynamics.
- Stocks like TSLA and AI had more **uncertain forecasts**, consistent with their high volatility.

Model Strengths:

- **Strength:** Captures linear trends and seasonality well.
- **Limitation:** Doesn't adapt to nonlinear changes or sudden shocks.

Inference:

ARIMA is a **reliable baseline model**, especially for steady-growth stocks. For stocks with complex patterns or nonlinearity (like AI or TSLA), ARIMA forecasts may oversimplify the dynamics.

## 2. LSTM Forecasting

Performance Highlights:

- LSTM models demonstrated **better trend tracking** for volatile stocks (e.g., **TSLA, AI**) compared to ARIMA.
- Visually, the LSTM predictions aligned more **smoothly with actual prices** for some stocks, though not all.
- Predicted vs. Actual plots showed a **lag** in some cases, due to limited sequence length and smoothing.

Forecasting Evaluation Metrics (LSTM Results)

Stock	MAE	RMSE	MAPE (%)	Adjusted SMAPE (%)
AI	1.60	2.09	5.84	7.3
AMD	7.89	10.18	6.32	8.1
AMZN	16.12	17.65	12.03	17.5
ANET	110.79	120.17	52.07	74.2
META	101.58	127.02	27.54	33.6
NOW	96.14	119.10	15.28	22.1
NVDA	3.94	5.63	7.28	9.5
PANW	154.91	179.56	58.36	81.4
PATH	4.18	4.47	18.67	25.2
TSLA	36.05	39.72	17.57	24.9

## Inference

- **Strong LSTM Performance:**
  - AI, AMD, NVDA – lowest MAE, RMSE, and <10% MAPE/SMAPE — **very stable and predictable.**
- **Moderate Forecast Accuracy:**
  - AMZN, NOW, TSLA, and PATH show decent forecasting, with **MAPE and SMAPE <30%**, making them **reasonably predictable.**
- **Low Forecast Accuracy:**
  - META, ANET, and PANW have **MAPE >25%** and **SMAPE >70%**, indicating **significant forecasting challenges**, likely due to volatility, sudden shifts, or irregular patterns.

### C. Volatility Analysis (*Rolling Standard Deviation*)

#### Volatility Patterns:

- TSLA and AI experienced **frequent spikes**, indicating rapid market reaction to events/news.
- META, NOW, and AMZN showed **lower rolling SD**, reflecting more **stable** price movement.
- Post-2022, a general **increase in volatility** was seen across most stocks—likely due to macroeconomic uncertainty.

#### Inference:

Volatility trends help investors **gauge risk**. Highly volatile stocks may offer high return potential but come with higher uncertainty. Rolling SD visualizations are vital for **portfolio diversification decisions**.